



IERG4210 Web Programming and Security

Course Website: <https://course.ie.cuhk.edu.hk/~ierg4210/>
Live FB Feedback Group: <https://fb.com/groups/ierg4210.2014spring/>

Fast and Scalable Web & Database Servers

Lecture 6

Dr. Adonis Fung
phfung@ie.cuhk.edu.hk

Information Engineering, CUHK
Product Security Engineering, Yahoo!

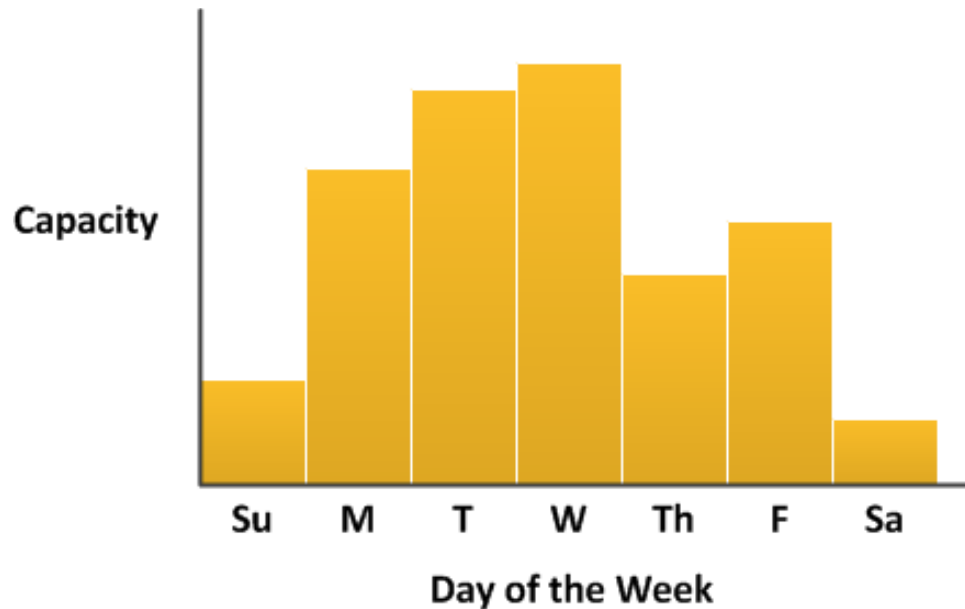
Agenda

- Fast and Scalable Web Servers
 - Quick Introduction to Cloud
 - Architecture and Designs
- DB Storage Servers
 - Quick Intro. to DB Storage
 - Database v.s. Cache
 - Relational Databases (MySQL, SQLite)
 - Structured Query Language (SQL) Language
 - Example Usage thru Database Abstraction Layer
 - Quick intro. to in-memory cache (Redis)

Why Cloud?

- **Uneven Utilizations by nature**

- Day of week and Hour of day: Web surfing
- Season of year: Christmas e-gift cards
- Adhoc usage: One-off computation jobs/testing



Cloud Benefits: on-demand + sharing

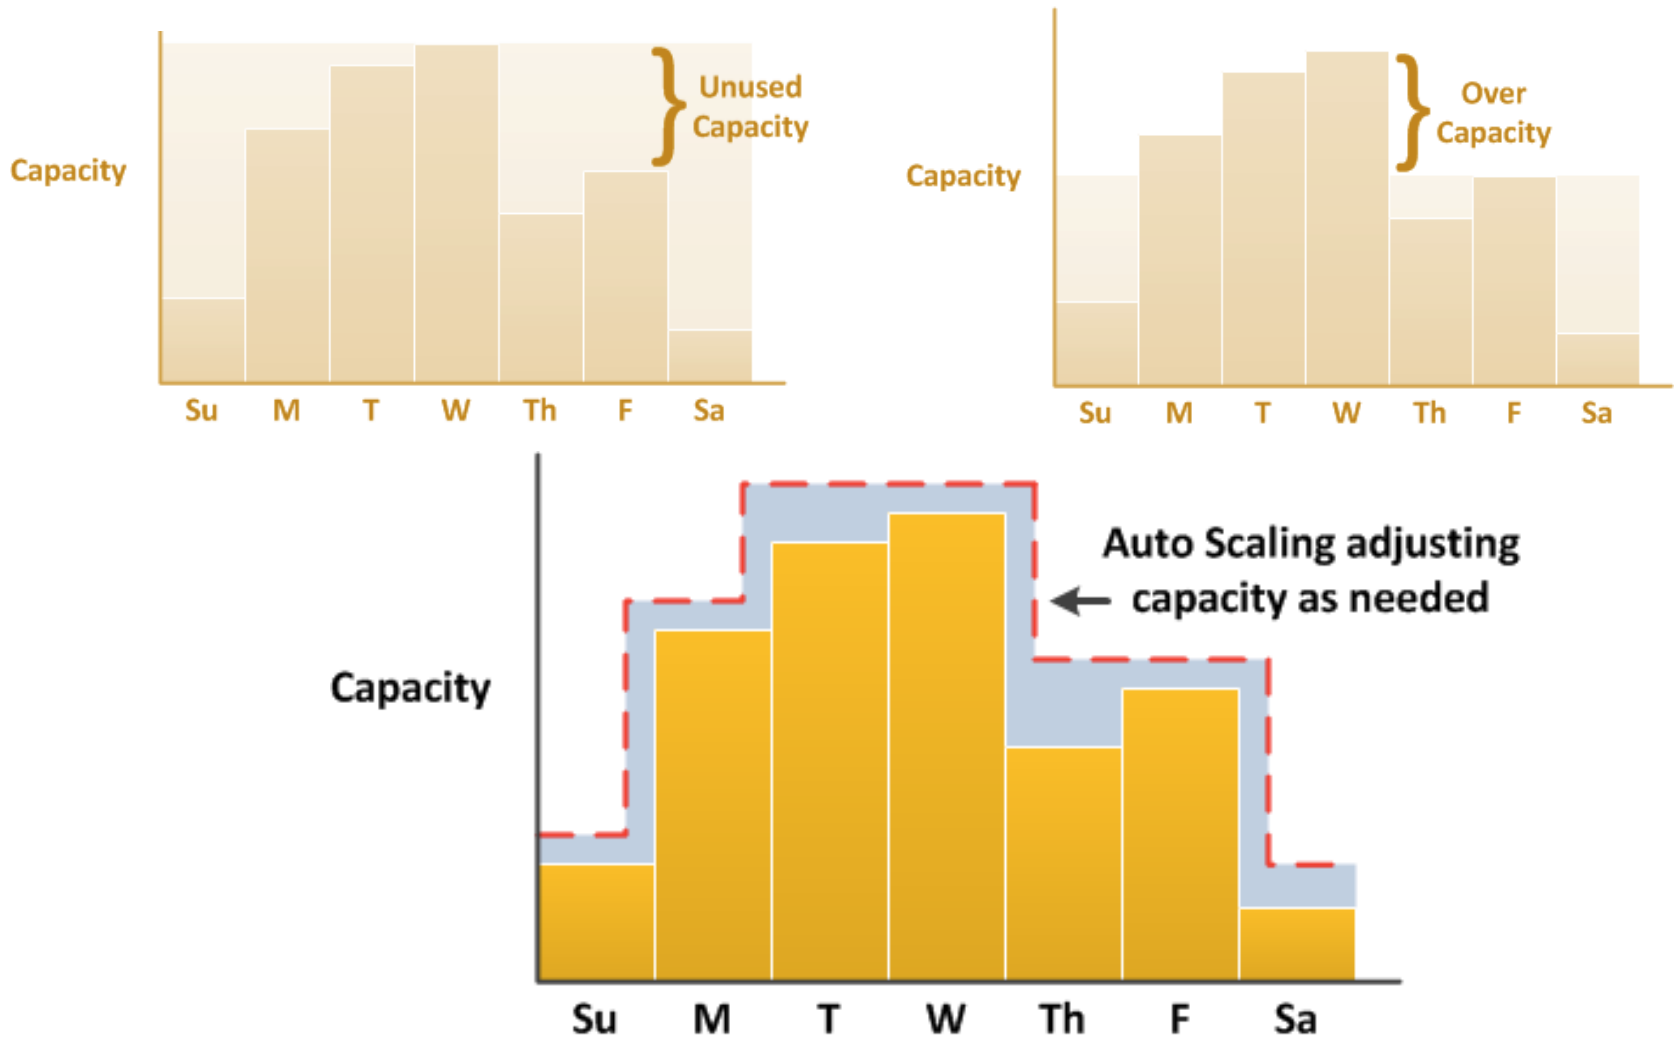


Image Ref: <http://docs.aws.amazon.com/AutoScaling/latest/DeveloperGuide/as-dg.pdf>

Cloud Classifications

- Cloud Service Models

- Infrastructure-as-a-Service (IaaS)
- Platform-as-a-Service (PaaS)
- Software-as-a-Service (SaaS)

Suddenly, everything become
X-as-a-Service (XaaS)

- What service models?

- AWS EC2
- AWS Elastic Beanstalk
- Google Cloud Console
- Google AppEngine

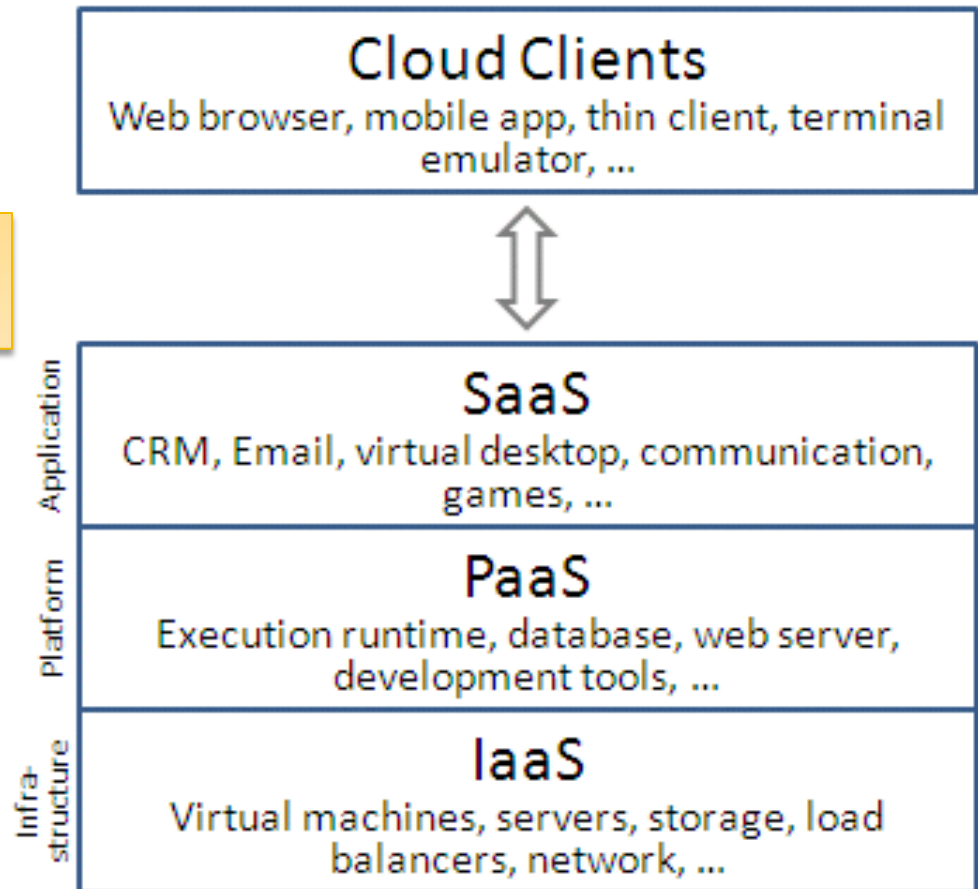
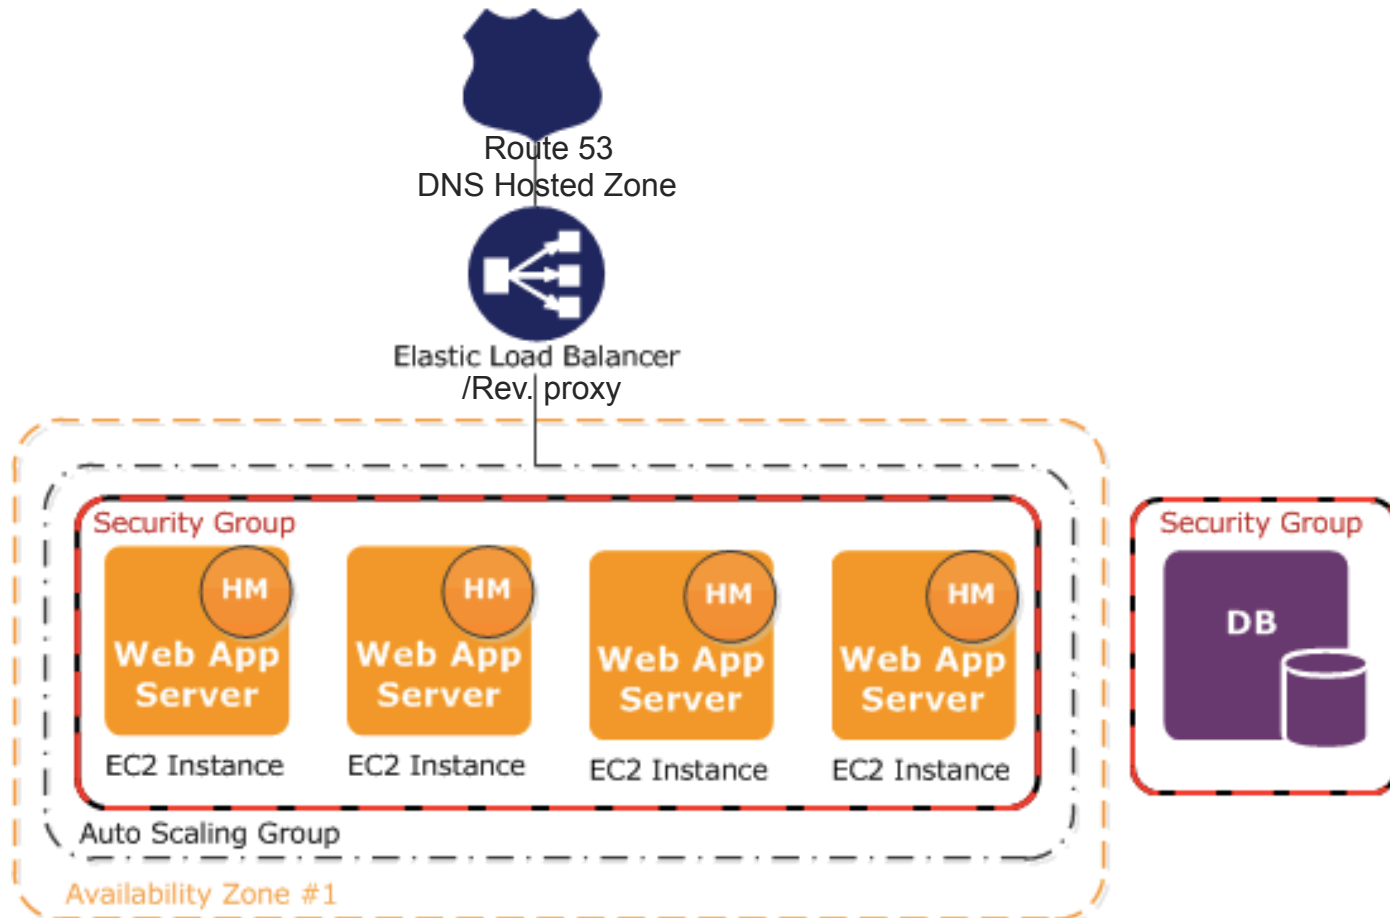


Image Ref: http://en.wikipedia.org/wiki/Cloud_computing#mediaviewer/File:Cloud_computing_layers.png

Fast and Scalable Web Server Platform

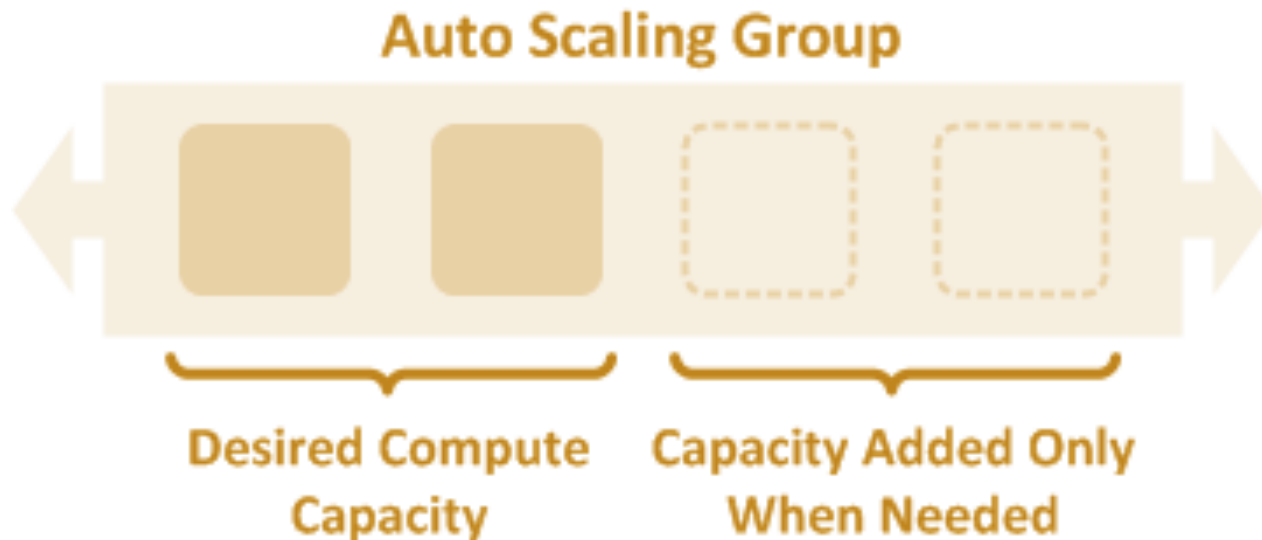
- **Architecture** of AWS Elastic Beanstalk:
Web Server Tiers



Details and Image Ref: <http://docs.aws.amazon.com/elasticbeanstalk/latest/dg/concepts.concepts.architecture.html>

Scale up v.s. Scale out

- **Scale up (vertically)**: more resources for a single node
 - More expensive for supercomputer (less efficient)
- **Scale out (horizontally)**: more nodes
 - A farm of cheaper instances well-networked (high throughput)
 - application distributable in idv. instance



Design Considerations

- Scalability and Elasticity
 - Auto Scaling enables on demand instance/node creation/removal
 - Possibly based on metrics: CPU, memory, disk I/O, network I/O, etc...
- Fault Tolerance for Availability
 - Automated recovery for EC2 instances when some of them die
 - Automated backups for Storage
- Software Deployability
 - Easy to deploy; and create new environment to test out changes
 - Integrated with GIT for systematic versioning control
- that's why we chose AWS EB
 - plus we're awarded the education grant :)

DB STORAGE SERVERS

Introduction to DB Storage

- Relational Databases

- Structured in tables : Slow but powerful
- Hard to scale
- Accessible through the Structured Query Language (SQL)
- Often used as persistence storage
- Examples: MySQL (free), MSSQL, Oracle, SQLite (free), etc

- NoSQL Databases

- Unstructured as a tradeoff for speed
- Easy to scale out, mostly query, async, inaccuracy tolerance by app
- Accessible through API
- Optimized for speed, thus often as In-memory Cache
- Examples: Redis (free), Memcached (free), MongoDB, etc

Database: SQLite



- Public domain license (i.e. FREE!)
- **Lightweight** in design
 - Lightweight: multiple processes can read at the same time; however, only one process can make changes at any moment in time
 - Best for **single-user** apps (MobileApps/Simple WebApps)
- Supported by **multi-platforms** (e.g. Windows, Linux)
 - Pre-installed in AWS EC2
- Stores everything in a **single file**
 - Easy to embed, test, backup and transfer
- **Simple access-right management**
 - No user account management as in full-blown DBs like MySQL
 - Simply depends on the file access rights

Database: MySQL



- Dual-licensing: GPL/FLOSS + proprietary
- Relational DB
 - Table structure
- Full-featured, accessible using SQL
 - But heavyweight, quite slow
 - Powerful as something for free
- Supported by multi-platforms (e.g. Windows, Linux)
 - Pre-installed in AWS EC2

MySQL favor

SQL LANGUAGE

Create a MySQL DB for EB

1. (local-env) \$ eb console
2. Click Configuration -> Under Data Tier, Click Create DB
3. [shop123-ierg4210](#) ▶ [shop123-ierg4210-dev](#) ([shop123-ierg4210-dev.elasticbeanstalk.com](#))

Dashboard	Snapshot	<input type="text" value="None"/>	Refresh
Configuration	DB Engine	<input type="text" value="mysql"/>	Refresh
Logs	Instance Class	<input type="text" value="db.t1.micro"/>	Refresh
Monitoring	Allocated Storage	<input type="text" value="5"/> <input type="text" value="GB"/>	You must specify a value in the range 5 GB to 1024 GB.
Alarms	Master Username	<input type="text" value="root"/>	
Events	Master Password	<input type="password" value="....."/>	
Tags	Deletion Policy	<input type="text" value="Create Snapshot"/>	Terminating your environment can permanently delete your Amazon RDS DB instance and all its data. a snapshot, which preserves your data but may incur backup storage charges. Learn more.
	Availability	<input type="text" value="Single Availability Zone"/>	

Connect to the DB using MySQL CLI

1. (local-env) \$ eb ssh
 2. \$ sudo yum install mysql -y
 3. \$ mysql -u root -p -h <your-db>.rds.amazonaws.com
 - Enter your configured Master Password
-
- It fails when connecting directly to MySQL from local. Why?
 - The EC2 instances and MySQL are in the same Security Group, hence in the same Virtual Private Cloud (network)
 - From Security Group settings, expect only port 80 is opened
 - SSH is hosted at port 22, and is dynamically made accessible by eb ssh
 - MySQL is hosted at 3306, and thus cannot be accessed by public
 - Hence, access the DB thru EC2

MySQL: Create a User and DB

Welcome to the MySQL monitor. Commands end with ; or \g.
Your MySQL connection id is 26
Server version: 5.5.40-log Source distribution

Copyright (c) 2000, 2014, Oracle and/or its affiliates. All rights reserved.

Oracle is a registered trademark of Oracle Corporation and/or its affiliates. Other names may be trademarks of their respective owners.

Type 'help;' or '\h' for help. Type '\c' to clear the current input statement.

```
mysql> CREATE USER 'shopXX-admin' IDENTIFIED BY 'mypass';  
mysql> CREATE DATABASE shopXX;  
mysql> SHOW DATABASES;  
mysql> GRANT ALL ON shopXX.* TO 'shopXX-admin';  
mysql> exit;
```


MySQL: Login using New User and Pick new DB

- Re-login using the newly created user
 - `$ mysql -u shopXX-admin -p -h <your-db>.rds.amazonaws.com`
- Pick the newly created DB
 - `mysql> USE shopXX;`
- Considerations:
 - shopXX-admin is granted full access to DB called shopXX
 - You may like practicing the least privilege approach (e.g., SELECT)
 - To drop/delete the user:
`mysql> DROP USER shopXX-admin`
 - Using shopXX-admin@localhost will fail because we'd support remote logins from the EC2

Create a Table (1/2)

- Creating a Table “categories”

- Create the table

```
mysql> CREATE TABLE categories (  
    catid INTEGER PRIMARY KEY AUTO_INCREMENT,  
    name VARCHAR(512) NOT NULL  
    ) ENGINE=INNODB;
```

Note: Primary key is unique and auto-increment by default
(i.e. incremented by 1 automatically for every new record)

- To check what you have created

```
mysql> DESCRIBE categories;
```

- To drop/delete the whole table and data

```
mysql> DROP TABLE categories;
```

- To drop/delete all the data

```
mysql> TRUNCATE categories;
```

Create a Table (2/2)

- Creating a Table “products”

- Create the table (simplified, add price/description type yourself)

```
mysql> CREATE TABLE products (  
    pid INTEGER PRIMARY KEY AUTO_INCREMENT,  
    catid INTEGER,  
    name VARCHAR(512),  
    price _____,  
    description _____,  
    FOREIGN KEY(catid)  
        REFERENCES categories(catid)  
    ) ENGINE=INNODB;
```

Reference: [Datatypes supported by MySQL](#)

- Create an index for catid – to make subsequent queries by catid faster

```
mysql> CREATE INDEX i1 ON products(catid);
```

INSERT

- Inserting some records to the newly created tables

- To insert a record into categories

```
mysql> INSERT INTO categories VALUES (null, "Fruits");
```

Note: put null for the primary key to let it auto-increment

- To insert 2 records into products

```
mysql> INSERT INTO products  
VALUES (null, 1, "Apple", "1.5");
```

```
mysql> INSERT INTO products (catid, name, price)  
VALUES (1, "Banana", "1.5");
```

- Try to insert a product to an inexistent category:

```
mysql> INSERT INTO products (catid, name, price)  
VALUES (2, "Help", "999");
```

Error: constraint failed

Note: This error is expected given that the foreign key setting

- More on SQL INSERT: <http://dev.mysql.com/doc/en/insert.html>

SELECT

- Looking up records

- To select all “fruits” in products (given fruits is of catid=1):

```
mysql> SELECT * FROM products WHERE catid = 1;
```

- To select only the name and price columns

```
mysql> SELECT name, price FROM products WHERE catid = 1;
```

- To select only 5 “fruits” in products:

```
mysql> SELECT * FROM products WHERE catid = 1 LIMIT 5;
```

- To select the 11-20th most expensive “fruits” in products: (Pagination?)

```
mysql> SELECT * FROM products  
        WHERE catid = 1  
        ORDER BY price DESC  
        LIMIT 11, 10;
```

- Recall: For those columns that are frequently queried, remember to create INDEX for performance sake (trading off space for speed)
- More on SQL SELECT: <http://dev.mysql.com/doc/en/select.html>

UPDATE

- Updating a record

- Setting a static value

```
mysql> UPDATE categories  
      SET name = "Fresh Fruits"  
      WHERE catid = 1;
```

- Setting an expression (e.g. 10% increase in price)

```
mysql> UPDATE products  
      SET price = price * 1.1  
      WHERE pid = 2;
```

- More on SQL UPDATE: <http://dev.mysql.com/doc/en/update.html>

- Note:

- The WHERE conditions is the same as that of SELECT
 - So, when you are not sure about what records are affected
 - SELECT the records first, then replace it with UPDATE
 - Otherwise, you can kill all your data unintentionally

DELETE

- Deleting a record
 - Deleting the fruit category will result in an error
`mysql> DELETE FROM categories WHERE catid = 1;`
Reason: Given foreign key is ON, a cat with children can't be deleted
 - Deleting a product
`mysql> DELETE FROM products WHERE pid = 2;`
 - The where conditions are again the same as that of SELECT
 - More on SQL DELETE: <http://dev.mysql.com/doc/en/delete.html>
- Deleting records requires extra attention!
 - Backup your database
 - Or SELECT what rows are affected before performing DELETE

Database Abstraction Layer

- A universal interface for accessing to different databases
 - **Coding Consistency**: Regardless of the DB, use the same set of code
 - **Single Interface**: Easy to switch database without code modifications




Image Ref: W. Jason Gilmore, Beginning PHP and MySQL From Novice to Professional, p.794, 2007

Example Usage

- Using [Any-DB](#),

```
var anyDB = require('any-db');
var config = require('../shopXX-ierg4210.config.js');
var pool = anyDB.createPool(config.mysqlURI, {
    min: 2, max: 20
});

app.get('/', function (req, res) {
    // async fetch data from SQL, render page when ready
    pool.query('SELECT * FROM categories', function (error, results) {
        if (error) {
            console.error(error);
            return res.status(500).end();
        }
        res.render('home', {
            title: 'IERG4210 ShopXX',
            cat: results.rows
        });
    });
});
```

Redis



- Open-source NoSQL DB/Cache
 - In-memory key-value store (hence, very fast)
 - But also supports data structures such as sorted sets and lists
- Common use cases:
 - To serve queries, therefore GET requests
e.g., cache your templates to prevent from re-rendering
 - When to expire? Expire on DB update?
 - To completely serve as a DB
 - Data loss when machine powers down (due to in memory)
 - Periodically backup data to persistent storage
- See [Redis NPM](#), [Redis.io](#), and [AWS ElastiCache](#) for details

Some Logistics...

- Assignment Phase 3 Released